

# Design Document

**System Objectives:** Build an AI-assisted coding practice platform that generates a complete problem-solving environment (formal specification, test suite, reference solutions, and an IDE) from minimal or vague user input. We understand a candidate after giving an interview “n” days ago may not be able to recall the problem fully well, so they can give a vague statement to our platform, and we will do our best to generate a coding environment for it. Whether a user pastes a known LeetCode statement (also called a formal specification) or recalls a problem in their own words, the system formalizes it, creates diverse tests, and provides a reference solution for practice. The platform emphasizes **robustness** and **trust**: any errors (in the spec, tests, or solution) can be flagged by the user with counterexamples, triggering automated self-healing to correct mistakes.

**User Personas:** We model four input “personas” reflecting how different people describe problems after an interview:

- *Layman*: Simple, non-technical phrasing. “Interview recall without terminology”. (e.g. “find the common string”).
- *Conversational*: Normal speech with some detail. “Human-to-human explanation”.
- *Technical Shorthand*: Shorter, jargon-heavy descriptions. “Abbreviated interview notes”.
- *Implementation-Specific*: Emphasis on code/pseudocode hints.

These personas simulate real users who might recall a problem vaguely. By fine-tuning on examples of these four styles, the platform can parse any user query without requiring formal math language.

**Design Principles:** We follow **completeness**, **correctness**, and **self-healing**. The system is **user-centric** (handles incomplete input), **reliable** (multiple verification layers), and **iterative** (allowing feedback loops). Key practices include:

- **Multi-stage pipeline**: Each problem goes through distinct stages (specification, test generation, solution generation, evaluation) to modularize complexity.
  - Stage 1: Formal Specification to Vague Statement (Training)
  - Stage 2: Fine-tuning vague and formal specification (Training)
  - Stage 3: Vague Problem Statement to Formal Specification (Inference)
  - Stage 4: Test Suite Generation (Inference)
  - Stage 5: Reference Solution Generation (Inference)
  - Stage 6: Candidate Solution Validation.
  - Stage 7: Candidate Feedback with Counterexamples.
- **Formal specifications**: Converting to a precise spec ensures clarity and correctness, following strict formatting rules (e.g. function signature

`solve(self, ...)`, sample examples and constraint extraction).

- **Automated test “oracle”**: LLM-generated test cases are validated by a local `is_valid_input` function (oracle) to ensure they respect constraints.
- **Brute-force baseline**: Brute-force solutions serve as ground truth for correctness; optimized solutions are compared against them (healing) to detect and fix errors.
- **Self-healing/cross-healing**: If tests or solutions fail (either fails Oracle or conflict between brute/optimized), the system automatically retries with adjusted prompts or logic until consistency is achieved (circuit-breaker strategy using max retries)
- **User feedback loop**: End users can provide counterexamples; a “Gatekeeper” audits them against constraints and, if valid, injects these into the pipeline to regenerate tests/solutions (truth-injection healing).
- **Scalability & efficiency**: The design assumes large DSA inputs; for stress tests it uses sparse representations (e.g. Python repeat syntax) and timeouts to prevent resource overuse.

**Fine-Tuning Rationale:** We curated **215 LeetCode problems** (various categories and difficulties) from a public dataset. For each, we took the formal LeetCode statement and generated four “vague” variants (one per persona) in English. This produced (informal description → formal specification) training pairs. We then fine-tuned a GPT-4.1-mini model on this dataset using supervised fine-tuning (SFT). The result is a specialized LLM that can translate user queries in any of the four styles into a precise problem spec (JSON with title, description, constraints, starter code, etc.). By anchoring on real LeetCode problems and these personas, the model learns to handle ambiguous or paraphrased inputs robustly.

**Training & Translation Strategy:** The HuggingFace LeetCode dataset provided ground-truth problem specs. For each, we wrote four persona-flavored prompts (e.g. an interviewee’s layman description) as input, and the original spec as output. During inference, the fine-tuned model uses a strict prompt (e.g. “Convert this vague question into a formal spec”) to generate JSON output. This supervised approach ensures alignment between informal recollection and formal problem definition.